RUPESH GHIMIRE
Milwaukee, WI 53211
+1 (469) 690-7056
rupacegh@gmail.com | Rupesh Ghimire - LinkedIn

PROFESSIONAL SUMMARY:

- Senior Data Engineer with 8 years of experience designing and building scalable data pipelines, ETL solutions, and data platforms for large-scale data processing.
- Strong hands-on experience with **Databricks**, including **Unity Catalog, Databricks Asset Bundles (DABs), Auto Loader, and PySpark** for modern data engineering and pipeline development.
- Proficient in **Python, SQL, PySpark, and Shell Scripting**, with a strong understanding of distributed data processing and Spark frameworks.
- Experienced in designing and optimizing **ETL/ELT pipelines** with a focus on data quality, performance, reliability, and maintainability.
- Delivered cloud and data migration solutions across **Databricks, AWS, and Azure**, including large-scale data processing workloads handling terabytes of data.
- Hands-on experience with **AWS data services** including S3, Glue, Athena, Lambda, AWS Batch, EC2, DynamoDB, and SQS/SNS, along with Azure Data Factory and Azure Storage.
- Experienced across modern and traditional data platforms, including **Databricks, Snowflake, Oracle**, and enterprise ETL technologies such as Informatica, ODI, Matillion, and Airflow.
- Proven experience leading development activities, coordinating QA, onboarding team members, working with cross-functional teams, and driving technical proof-of-concepts from evaluation through implementation.
- Domain experience delivering data and analytics solutions across **Insurance, Finance, and Retail**, with an emphasis on scalable cloud data platforms and **AI-assisted development workflows** for faster coding, debugging, refactoring, and solution delivery.

TECHNICAL SKILLS:

Big Data	Databricks, Apache Spark, PySpark,
ETL Tools	Informatica PowerCenter, ODI, Matillion, Apache Air Flow
Programming Languages	Python, SQL, PL/SQL, Shell Scripting
Cloud Platform	Azure, AWS, AWS Services (S3, AWS Glue, Athena, Lambda, Batch, EC2)
Data Warehousing Platforms	Databricks, Snowflake, Oracle
OLAP BI Analytics	Power BI, Looker, OBIEE
Version control	GIT
Project Management	Agile, Scrum, Atlassian, GitLab, JIRA

EDUCATION:

- Bachelor of Engineering in Computer Engineering, 2014 - 2018
  Kathmandu University – Nepal

WORK HISTORY:

Northwestern Mutual | Milwaukee, WI | Sep 2023 – Present
Senior Software Engineer

Responsibilities:

- Design and develop **enterprise-scale data pipelines on Databricks** using Python, PySpark, SQL, and Delta Lake, loading curated data into enterprise data tables supporting Insurance, Client, Wealth, and Field Advisor domains.
- Develop and maintain data solutions using **Databricks and Unity Catalog**, with a focus on scalable architecture, data governance, reliability, and secure access to enterprise data.
- Build reusable **ETL frameworks, libraries, and common data ingestion components** that standardize pipeline development and are adopted across engineering teams to improve consistency and development efficiency.
- Partner closely with **business teams, product managers, and engineering leadership** to translate business requirements into scalable data pipelines and reliable analytical data products.
- Design and support multiple **data ingestion patterns**, including Kafka-based ingestion, Databricks pipelines, Control-M job orchestration, and automated data validation workflows.
- Implement and support **data quality and validation processes**, including Databricks DQX, reconciliation, and automated checks to improve the accuracy and reliability of enterprise data.
- Drive **pipeline performance and cost optimization**, redesigning Databricks workloads and engineering solutions to reduce job costs, in some cases by up to **70% per run** while improving runtime by up to **30%**.
- Lead technical **proofs of concept and solution evaluations** across Databricks and other technologies to identify practical approaches for improving scalability, performance, cost, and developer productivity.
- Explore and implement **AI-assisted development and AI agents use cases** for data engineering, including custom skills, instructions, and workflows that help developers automate repetitive tasks and accelerate pipeline development.
- Contribute as a senior engineer through **technical design, framework development, code reviews, troubleshooting, knowledge sharing, and mentoring**, helping establish engineering standards across the team.

Environment: Databricks, Unity Catalog, PySpark, Python, SQL, Delta Lake, Databricks DQX, Kafka, Control-M, ETL/ELT, Data Ingestion, Git, Azure, AI/AI Agents

Verisk Analytics | Jersey City, NJ (Remote) | Feb 2021 – Feb 2023
Senior Software Engineer

Responsibilities:

- Designed and developed Python and Shell-based ETL pipelines for large-scale data processing and data warehousing on AWS.
- Designed and implemented end-to-end data pipelines processing terabytes of data per day using AWS services, PySpark, Python, and SQL.
- Developed and optimized data models and complex SQL/ETL processes to support analytical and reporting requirements.
- Built and maintained Snowflake data warehouse solutions, including data structures and analytics-ready datasets.
- Automated data loading, partition creation, and file conversion using Python, AWS Lambda, and AWS Batch, reducing manual effort by approximately 70%.
- Built configuration-driven data processing solutions that support multiple clients and use cases through reusable JSON-based configurations.

- Worked extensively with flat-file data ingestion and extraction, integrating file-based data with cloud data platforms.
- Collaborated with data scientists, software engineers, and business stakeholders to develop reliable data solutions and improve data-driven business outcomes.
- Supported ETL development and production stability by validating data integrity, pipeline reliability, and processing performance.

Environment: Python, AWS Services, PySpark, SQL, Snowflake, Shell Scripting, Athena, AWS Batch, JSON

Logic Information System Inc – Minneapolis, Minnesota (Remote) | Sept 2018 – Feb 2021
Software Engineer

Responsibilities:

- Designed and developed ETL pipelines and data integration solutions using Python, Apache Airflow, Azure, and SQL, processing terabytes of data per week.
- Built and maintained data warehouse solutions for ERP data, integrating order management, invoicing, and service contract data for reporting and analytics.
- Developed complex Oracle Data Integrator (ODI) ETL workflows to migrate and transform data from flat files into relational databases.
- Designed scalable Snowflake data warehouse architectures, including tables, views, materialized views, and data integration pipelines optimized for analytical workloads.
- Integrated data from multiple source systems using ETL tools and custom data pipelines, ensuring reliable and consistent data delivery.
- Improved data processing performance by applying Apache Spark and distributed computing techniques and optimizing SQL queries.
- Developed Shell scripts to automate file transfers, notifications, and other data processing and operational tasks.
- Built Power BI and OBIEE dashboards and developed ad hoc Oracle reports to support business analysis and performance reporting.
- Performed data validation and reconciliation between source systems and data warehouses, identifying and resolving data quality and business logic issues.
- Implemented data security controls, including encryption, access control, and monitoring, to support data privacy and compliance requirements.

Environment: Python, Shell Script, Oracle, SQL, Snowflake, Azure, Power BI